

Understanding JSON Schema

Release 7.0

Michael Droettboom, et al
Space Telescope Science Institute

Jul 10, 2021

Contents

1 Conventions used in this book	3
1.1 Language-specific notes	3
1.2 Draft-specific notes	4
1.3 Examples	4
2 What is a schema?	7
3 The basics	11
3.1 Hello, World!	11
3.2 The type keyword	12
3.3 Declaring a JSON Schema	13
3.4 Declaring a unique identifier	13
4 JSON Schema Reference	15
4.1 Type-specific keywords	15
4.2 string	17
4.2.1 Length	19
4.2.2 Regular Expressions	19
4.2.3 Format	20
4.3 Regular Expressions	22
4.3.1 Example	23
4.4 Numeric types	23
4.4.1 integer	24
4.4.2 number	25
4.4.3 Multiples	25
4.4.4 Range	26
4.5 object	29
4.5.1 Properties	30
4.5.2 Pattern Properties	31
4.5.3 Additional Properties	32
4.5.4 Required Properties	34
4.5.5 Property names	36
4.5.6 Size	36
4.6 array	37
4.6.1 Items	38
4.6.2 Additional Items	40
4.6.3 Contains	42

4.6.4	Length	43
4.6.5	Uniqueness	44
4.7	boolean	44
4.8	null	45
4.9	Generic keywords	46
4.9.1	Annotations	46
4.9.2	Comments	47
4.9.3	Enumerated values	47
4.9.4	Constant values	48
4.10	Media: string-encoding non-JSON data	49
4.10.1	contentMediaType	49
4.10.2	contentEncoding	49
4.10.3	Examples	49
4.11	Schema Composition	50
4.11.1	allOf	51
4.11.2	anyOf	51
4.11.3	oneOf	52
4.11.4	not	53
4.11.5	Properties of Schema Composition	53
4.12	Applying subschemas conditionally	56
4.12.1	Dependencies	57
4.12.2	If-Then-Else	61
4.12.3	Implication	65
4.13	Declaring a Specification Version	66
4.13.1	\$schema	66
4.13.2	Extending JSON Schema	67
5	Structuring a complex schema	71
5.1	Schema Identification	71
5.1.1	JSON Pointer	72
5.1.2	Named Anchors	72
5.2	Base URI	73
5.2.1	Retrieval URI	73
5.2.2	\$id	74
5.3	\$ref	75
5.4	definitions	76
5.5	Recursion	77
5.6	Bundling	78
6	Acknowledgments	81
Index		83

JSON Schema is a powerful tool for validating the structure of JSON data. However, learning to use it by reading its specification is like learning to drive a car by looking at its blueprints. You don't need to know how an electric motor fits together if all you want to do is pick up the groceries. This book, therefore, aims to be the friendly driving instructor for JSON Schema. It's for those that want to write it and understand it, but maybe aren't interested in building their own car—er, writing their own JSON Schema validator—just yet.

Note: This book describes JSON Schema draft 7. Earlier versions of JSON Schema are not completely compatible with the format described here, but for the most part, those differences are noted in the text.

Where to begin?

- This book uses some novel *conventions* (page 3) for showing schema examples and relating JSON Schema to your programming language of choice.
- If you're not sure what a schema is, check out *What is a schema?* (page 7).
- *The basics* (page 11) chapter should be enough to get you started with understanding the core *JSON Schema Reference* (page 15).
- When you start developing large schemas with many nested and repeated sections, check out *Structuring a complex schema* (page 71).
- [json-schema.org](#) has a number of resources, including the official specification and tools for working with JSON Schema from various programming languages.
- There are a number of [online JSON Schema tools](#) that allow you to run your own JSON schemas against example documents. These can be very handy if you want to try things out without installing any software.

CHAPTER 1

Conventions used in this book

- *Language-specific notes* (page 3)
- *Draft-specific notes* (page 4)
- *Examples* (page 4)

1.1 Language-specific notes

The names of the basic types in JavaScript and JSON can be confusing when coming from another dynamic language. I'm a Python programmer by day, so I've noted here when the names for things are different from what they are in Python, and any other Python-specific advice for using JSON and JSON Schema. I'm by no means trying to create a Python bias to this book, but it is what I know, so I've started there. In the long run, I hope this book will be useful to programmers of all stripes, so if you're interested in translating the Python references into Algol-68 or any other language you may know, pull requests are welcome!

The language-specific sections are shown with tabs for each language. Once you choose a language, that choice will be remembered as you read on from page to page.

For example, here's a language-specific section with advice on using JSON in a few different languages:

Python

In Python, JSON can be read using the json module in the standard library.

Ruby

In Ruby, JSON can be read using the json gem.

C

For C, you may want to consider using [Jansson](#) to read and write JSON.

1.2 Draft-specific notes

The JSON Schema standard has been through a number of revisions or “drafts”. The current version is Draft 7, but Draft 4 is still widely used as well.

The text is written to encourage the use of Draft 7 and gives priority to the latest conventions and features, but where it differs from earlier drafts, those differences are highlighted in special call-outs. If you only wish to target Draft 7, you can safely ignore those sections.

New in draft 7

Draft 4

This is where anything pertaining to an old draft would be mentioned.

1.3 Examples

There are many examples throughout this book, and they all follow the same format. At the beginning of each example is a short JSON schema, illustrating a particular principle, followed by short JSON snippets that are either valid or invalid against that schema. Valid examples are in green, with a checkmark. Invalid examples are in red, with a cross. Often there are comments in between to explain why something is or isn’t valid.

Note: These examples are tested automatically whenever the book is built, so hopefully they are not just helpful, but also correct!

For example, here’s a snippet illustrating how to use the number type:

{ json schema }

```
{ "type": "number" }
```



42



-1

Simple floating point number:



5.0

Exponential notation also works:



2.99792458e8

Numbers as strings are rejected:



"42"

CHAPTER 2

What is a schema?

If you've ever used XML Schema, RelaxNG or ASN.1 you probably already know what a schema is and you can happily skip along to the next section. If all that sounds like gobbledegook to you, you've come to the right place. To define what JSON Schema is, we should probably first define what JSON is.

JSON stands for “JavaScript Object Notation”, a simple data interchange format. It began as a notation for the world wide web. Since JavaScript exists in most web browsers, and JSON is based on JavaScript, it's very easy to support there. However, it has proven useful enough and simple enough that it is now used in many other contexts that don't involve web surfing.

At its heart, JSON is built on the following data structures:

- object:

```
{ "key1": "value1", "key2": "value2" }
```

- array:

```
[ "first", "second", "third" ]
```

- number:

```
42  
3.1415926
```

- string:

```
"This is a string"
```

- boolean:

```
true  
false
```

- null:

null

These types have analogs in most programming languages, though they may go by different names.

Python

The following table maps from the names of JSON types to their analogous types in Python:

JSON	Python
string	string
number	int/float
object	dict
array	list
boolean	bool
null	None

45

⁴ Since JSON strings always support unicode, they are analogous to unicode on Python 2.x and str on Python 3.x.

⁵ JSON does not have separate types for integer and floating-point.

Ruby

The following table maps from the names of JSON types to their analogous types in Ruby:

JSON	Ruby
string	String
number	Integer/Float
object	Hash
array	Array
boolean	TrueClass/FalseClass
null	NilClass

6

⁶ JSON does not have separate types for integer and floating-point.

With these simple data types, all kinds of structured data can be represented. With that great flexibility comes great responsibility, however, as the same concept could be represented in myriad ways. For example, you could imagine representing information about a person in JSON in different ways:

```
{  
  "name": "George Washington",  
  "birthday": "February 22, 1732",  
  "address": "Mount Vernon, Virginia, United States"  
}  
  
{  
  "first_name": "George",  
  "last_name": "Washington",  
  "middle_name": null,  
  "date_of_birth": "1732-02-22",  
  "place_of_birth": "Mount Vernon, Virginia, United States",  
  "titles": ["First President of the United States"]  
}
```

(continues on next page)

(continued from previous page)

```

"last_name": "Washington",
"birthday": "1732-02-22",
"address": {
  "street_address": "3200 Mount Vernon Memorial Highway",
  "city": "Mount Vernon",
  "state": "Virginia",
  "country": "United States"
}
}

```

Both representations are equally valid, though one is clearly more formal than the other. The design of a record will largely depend on its intended use within the application, so there's no right or wrong answer here. However, when an application says “give me a JSON record for a person”, it's important to know exactly how that record should be organized. For example, we need to know what fields are expected, and how the values are represented. That's where JSON Schema comes in. The following JSON Schema fragment describes how the second example above is structured. Don't worry too much about the details for now. They are explained in subsequent chapters.

{ json schema }

```
{
  "type": "object",
  "properties": {
    "first_name": { "type": "string" },
    "last_name": { "type": "string" },
    "birthday": { "type": "string", "format": "date" },
    "address": {
      "type": "object",
      "properties": {
        "street_address": { "type": "string" },
        "city": { "type": "string" },
        "state": { "type": "string" },
        "country": { "type": "string" }
      }
    }
  }
}
```

By “validating” the first example against this schema, you can see that it fails:



```
{
  "name": "George Washington",
  "birthday": "February 22, 1732",
  "address": "Mount Vernon, Virginia, United States"
}
```

However, the second example passes:



```
{  
  "first_name": "George",  
  "last_name": "Washington",  
  "birthday": "1732-02-22",  
  "address": {  
    "street_address": "3200 Mount Vernon Memorial Highway",  
    "city": "Mount Vernon",  
    "state": "Virginia",  
    "country": "United States"  
  }  
}
```

You may have noticed that the JSON Schema itself is written in JSON. It is data itself, not a computer program. It's just a declarative format for “describing the structure of other data”. This is both its strength and its weakness (which it shares with other similar schema languages). It is easy to concisely describe the surface structure of data, and automate validating data against it. However, since a JSON Schema can't contain arbitrary code, there are certain constraints on the relationships between data elements that can't be expressed. Any “validation tool” for a sufficiently complex data format, therefore, will likely have two phases of validation: one at the schema (or structural) level, and one at the semantic level. The latter check will likely need to be implemented using a more general-purpose programming language.

CHAPTER 3

The basics

- *Hello, World!* (page 11)
- *The type keyword* (page 12)
- *Declaring a JSON Schema* (page 13)
- *Declaring a unique identifier* (page 13)

In *What is a schema?* (page 7), we described what a schema is, and hopefully justified the need for schema languages. Here, we proceed to write a simple JSON Schema.

3.1 Hello, World!

When learning any new language, it's often helpful to start with the simplest thing possible. In JSON Schema, an empty object is a completely valid schema that will accept any valid JSON.

{ json schema }

{ }

This accepts anything, as long as it's valid JSON



42



"I'm a string"



{ "an": ["arbitrarily", "nested"], "data": "structure" }

New in draft 6

You can also use `true` in place of the empty object to represent a schema that matches anything, or `false` for a schema that matches nothing.

{ json schema }

`true`

This accepts anything, as long as it's valid JSON



42



"I'm a string"



{ "an": ["arbitrarily", "nested"], "data": "structure" }

{ json schema }

`false`



"Resistance is futile... This will always fail!!!"

3.2 The type keyword

Of course, we wouldn't be using JSON Schema if we wanted to just accept any JSON document. The most common thing to do in a JSON Schema is to restrict to a specific type. The `type` keyword is used for that.

Note: When this book refers to JSON Schema “keywords”, it means the “key” part of the key/value pair in an object. Most of the work of writing a JSON Schema involves mapping a special “keyword” to a value within an object.

For example, in the following, only strings are accepted:

The diagram illustrates the validation of a JSON schema and its instances. At the top, a grey box labeled '{ json schema }' contains the schema: { "type": "string" }. Below it, a green box contains the valid instance "I'm a string" with a green checkmark icon. A red box contains the invalid instance 42 with a red X icon.

The type keyword is described in more detail in [Type-specific keywords](#) (page 15).

3.3 Declaring a JSON Schema

It's not always easy to tell which draft a JSON Schema is using. You can use the \$schema keyword to declare which version of the JSON Schema specification the schema is written to. See [\\$schema](#) (page 66) for more information. It's generally good practice to include it, though it is not required.

Note: For brevity, the \$schema keyword isn't included in most of the examples in this book, but it should always be used in the real world.

The diagram shows a JSON schema declaration: { "\$schema": "http://json-schema.org/draft-07/schema#" }. Below it, a black box labeled 'Draft 4' contains a note: 'In Draft 4, a \$schema value of http://json-schema.org/schema# referred to the latest version of JSON Schema. This usage has since been deprecated and the use of specific version URIs is required.'

3.4 Declaring a unique identifier

It is also best practice to include an \$id property as a unique identifier for each schema. For now, just set it to a URL at a domain you control, for example:

```
{ "$id": "http://yourdomain.com/schemas/myschema.json" }
```

The details of *\$id* (page 74) become more apparent when you start *Structuring a complex schema* (page 71).

New in draft 6

Draft 4

In Draft 4, *\$id* is just *id* (without the dollar-sign).

JSON Schema Reference

4.1 Type-specific keywords

The type keyword is fundamental to JSON Schema. It specifies the data type for a schema.

At its core, JSON Schema defines the following basic types:

- [*string*](#) (page 17)
- [*number*](#) (page 25)
- [*integer*](#) (page 24)
- [*object*](#) (page 29)
- [*array*](#) (page 37)
- [*boolean*](#) (page 44)
- [*null*](#) (page 45)

These types have analogs in most programming languages, though they may go by different names.

Python

The following table maps from the names of JSON types to their analogous types in Python:

JSON	Python
string	string
number	int/float
object	dict
array	list
boolean	bool
null	None

45

⁴ Since JSON strings always support unicode, they are analogous to `unicode` on Python 2.x and `str` on Python 3.x.

⁵ JSON does not have separate types for integer and floating-point.

Ruby

The following table maps from the names of JSON types to their analogous types in Ruby:

JSON	Ruby
string	String
number	Integer/Float
object	Hash
array	Array
boolean	TrueClass/FalseClass
null	NilClass

6

⁶ JSON does not have separate types for integer and floating-point.

The `type` keyword may either be a string or an array:

- If it's a string, it is the name of one of the basic types above.
- If it is an array, it must be an array of strings, where each string is the name of one of the basic types, and each element is unique. In this case, the JSON snippet is valid if it matches *any* of the given types.

Here is a simple example of using the `type` keyword:

{ json schema }

```
{ "type": "number" }
```



42



42.0

This is not a number, it is a string containing a number.



"42"

In the following example, we accept strings and numbers, but not structured data types:

{ json schema }

{ "type": ["number", "string"] }



42



"Life, the universe, and everything"



["Life", "the universe", "and everything"]

For each of these types, there are keywords that only apply to those types. For example, numeric types have a way of specifying a numeric range, that would not be applicable to other types. In this reference, these validation keywords are described along with each of their corresponding types in the following chapters.

4.2 string

- *Length* (page 19)
- *Regular Expressions* (page 19)
- *Format* (page 20)
 - *Built-in formats* (page 20)
 - * *Dates and times* (page 20)
 - * *Email addresses* (page 21)
 - * *Hostnames* (page 21)

- * *IP Addresses* (page 21)
- * *Resource identifiers* (page 21)
- * *URI template* (page 21)
- * *JSON Pointer* (page 22)
- * *Regular Expressions* (page 22)

The `string` type is used for strings of text. It may contain Unicode characters.

Python

In Python, "string" is analogous to the `unicode` type on Python 2.x, and the `str` type on Python 3.x.

Ruby

In Ruby, "string" is analogous to the `String` type.

{ json schema }

```
{ "type": "string" }
```



"This is a string"

Unicode characters:



"Déjà vu"



" "



"42"



42

4.2.1 Length

The length of a string can be constrained using the `minLength` and `maxLength` keywords. For both keywords, the value must be a non-negative number.

{ json schema }

```
{
  "type": "string",
  "minLength": 2,
  "maxLength": 3
}
```

	"A"
	"AB"
	"ABC"
	"ABCD"

4.2.2 Regular Expressions

The `pattern` keyword is used to restrict a string to a particular regular expression. The regular expression syntax is the one defined in JavaScript ([ECMA 262](#) specifically). See *Regular Expressions* (page 22) for more information.

Note: When defining the regular expressions, it's important to note that the string is considered valid if the expression matches anywhere within the string. For example, the regular expression "`p`" will match any string with a `p` in it, such as "apple" not just a string that is simply "p". Therefore, it is usually less confusing, as a matter of course, to surround the regular expression in `^...$`, for example, "`^p$`", unless there is a good reason not to do so.

The following example matches a simple North American telephone number with an optional area code:

{ json schema }

```
{  
  "type": "string",  
  "pattern": "^(\\([0-9]{3}\\))?[0-9]{3}-[0-9]{4}$"  
}
```



"555-1212"



"(888)555-1212"



"(888)555-1212 ext. 532"



"(800)FLOWERS"

4.2.3 Format

The `format` keyword allows for basic semantic validation on certain kinds of string values that are commonly used. This allows values to be constrained beyond what the other tools in JSON Schema, including [Regular Expressions](#) (page 22) can do.

Note: JSON Schema implementations are not required to implement this part of the specification, and many of them do not.

There is a bias toward networking-related formats in the JSON Schema specification, most likely due to its heritage in web technologies. However, custom formats may also be used, as long as the parties exchanging the JSON documents also exchange information about the custom format types. A JSON Schema validator will ignore any format type that it does not understand.

Built-in formats

The following is the list of formats specified in the JSON Schema specification.

Dates and times

Dates and times are represented in RFC 3339, section 5.6. This is a subset of the date format also commonly known as ISO8601 format.

- "date-time": Date and time together, for example, 2018-11-13T20:20:39+00:00.
- "time": New in draft 7 Time, for example, 20:20:39+00:00
- "date": New in draft 7 Date, for example, 2018-11-13.

Email addresses

- "email": Internet email address, see [RFC 5322](#), section 3.4.1.
- "idn-email": New in draft 7 The internationalized form of an Internet email address, see [RFC 6531](#).

Hostnames

- "hostname": Internet host name, see [RFC 1034](#), section 3.1.
- "idn-hostname": New in draft 7 An internationalized Internet host name, see [RFC5890](#), section 2.3.2.3.

IP Addresses

- "ipv4": IPv4 address, according to dotted-quad ABNF syntax as defined in [RFC 2673](#), section 3.2.
- "ipv6": IPv6 address, as defined in [RFC 2373](#), section 2.2.

Resource identifiers

- "uri": A universal resource identifier (URI), according to [RFC3986](#).
- "uri-reference": New in draft 6 A URI Reference (either a URI or a relative-reference), according to [RFC3986](#), section 4.1.
- "iri": New in draft 7 The internationalized equivalent of a "uri", according to [RFC3987](#).
- "iri-reference": New in draft 7 The internationalized equivalent of a "uri-reference", according to [RFC3987](#)

If the values in the schema have the ability to be relative to a particular source path (such as a link from a webpage), it is generally better practice to use "uri-reference" (or "iri-reference") rather than "uri" (or "iri"). "uri" should only be used when the path must be absolute.

Draft 4

Draft 4 only includes "uri", not "uri-reference". Therefore, there is some ambiguity around whether "uri" should accept relative paths.

URI template

- "uri-template": New in draft 6 A URI Template (of any level) according to [RFC6570](#). If you don't already know what a URI Template is, you probably don't need this value.

JSON Pointer

- "json-pointer": New in draft 6 A JSON Pointer, according to [RFC6901](#). There is more discussion on the use of JSON Pointer within JSON Schema in *Structuring a complex schema* (page 71). Note that this should be used only when the entire string contains only JSON Pointer content, e.g. /foo/bar. JSON Pointer URI fragments, e.g. #/foo/bar/ should use "uri-reference".
- "relative-json-pointer": New in draft 7 A [relative JSON pointer](#).

Regular Expressions

- "regex": New in draft 7 A regular expression, which should be valid according to the [ECMA 262](#) dialect.

Be careful, in practice, JSON schema validators are only required to accept the safe subset of *Regular Expressions* (page 22) described elsewhere in this document.

4.3 Regular Expressions

- *Example* (page 23)

The [pattern](#) (page 19) and [Pattern Properties](#) (page 31) keywords use regular expressions to express constraints. The regular expression syntax used is from JavaScript ([ECMA 262](#), specifically). However, that complete syntax is not widely supported, therefore it is recommended that you stick to the subset of that syntax described below.

- A single unicode character (other than the special characters below) matches itself.
- .: Matches any character except line break characters. (Be aware that what constitutes a line break character is somewhat dependent on your platform and language environment, but in practice this rarely matters).
- ^: Matches only at the beginning of the string.
- \$: Matches only at the end of the string.
- (...): Group a series of regular expressions into a single regular expression.
- |: Matches either the regular expression preceding or following the | symbol.
- [abc]: Matches any of the characters inside the square brackets.
- [a-z]: Matches the range of characters.
- [^abc]: Matches any character *not* listed.
- [^a-z]: Matches any character outside of the range.
- +: Matches one or more repetitions of the preceding regular expression.
- *: Matches zero or more repetitions of the preceding regular expression.
- ?: Matches zero or one repetitions of the preceding regular expression.
- +?, *?, ???: The *, +, and ? qualifiers are all greedy; they match as much text as possible. Sometimes this behavior isn't desired and you want to match as few characters as possible.
- (?!=x), (?=x): Negative and positive lookahead.
- {x}: Match exactly x occurrences of the preceding regular expression.
- {x,y}: Match at least x and at most y occurrences of the preceding regular expression.

- {x, }: Match x occurrences or more of the preceding regular expression.
- {x}?, {x,y}?, {x, }?: Lazy versions of the above expressions.

4.3.1 Example

The following example matches a simple North American telephone number with an optional area code:

```
{json schema}
{
  "type": "string",
  "pattern": "^(\\([0-9]{3}\\))?[0-9]{3}-[0-9]{4}$"
}
```

"555-1212"

"(888)555-1212"

"(888)555-1212 ext. 532"

"(800)FLOWERS"

4.4 Numeric types

- *integer* (page 24)
- *number* (page 25)
- *Multiples* (page 25)
- *Range* (page 26)

There are two numeric types in JSON Schema: *integer* (page 24) and *number* (page 25). They share the same validation keywords.

Note: JSON has no standard way to represent complex numbers, so there is no way to test for them in JSON Schema.

4.4.1 integer

The integer type is used for integral numbers. JSON does not have distinct types for integers and floating-point values. Therefore, the presence or absence of a decimal point is not enough to distinguish between integers and non-integers. For example, 1 and 1.0 are two ways to represent the same value in JSON. JSON Schema considers that value an integer no matter which representation was used.

Python

In Python, "integer" is analogous to the int type.

Ruby

In Ruby, "integer" is analogous to the Integer type.

{ json schema }

```
{ "type": "integer" }
```



42



-1

Numbers with a zero fractional part are considered integers



1.0

Floating point numbers are rejected:



3.1415926

Numbers as strings are rejected:



"42"

4.4.2 number

The `number` type is used for any numeric type, either integers or floating point numbers.

Python

In Python, "number" is analogous to the `float` type.

Ruby

In Ruby, "number" is analogous to the `Float` type.

{ json schema }

```
{ "type": "number" }
```



42



-1

Simple floating point number:



5.0

Exponential notation also works:



2.99792458e8

Numbers as strings are rejected:



"42"

4.4.3 Multiples

Numbers can be restricted to a multiple of a given number, using the `multipleOf` keyword. It may be set to any positive number.

{ json schema }

```
{  
  "type": "number",  
  "multipleOf" : 10  
}
```



0



10



20

Not a multiple of 10:



23

4.4.4 Range

Ranges of numbers are specified using a combination of the `minimum` and `maximum` keywords, (or `exclusiveMinimum` and `exclusiveMaximum` for expressing exclusive range).

If x is the value being validated, the following must hold true:

- $x \geq \text{minimum}$
- $x > \text{exclusiveMinimum}$
- $x \leq \text{maximum}$
- $x < \text{exclusiveMaximum}$

While you can specify both of `minimum` and `exclusiveMinimum` or both of `maximum` and `exclusiveMaximum`, it doesn't really make sense to do so.

{ json schema }

```
{  
  "type": "number",  
  "minimum": 0,  
  "exclusiveMaximum": 100  
}
```

Less than `minimum`:

	-1

`minimum` is inclusive, so 0 is valid:

	0

	10

	99

`exclusiveMaximum` is exclusive, so 100 is not valid:

	100

Greater than `maximum`:

	101

Draft 4

In JSON Schema Draft 4, `exclusiveMinimum` and `exclusiveMaximum` work differently. There they are boolean values, that indicate whether `minimum` and `maximum` are exclusive of the value. For example:

- if `exclusiveMinimum` is `false`, $x \geq \text{minimum}$.
- if `exclusiveMinimum` is `true`, $x > \text{minimum}$.

This was changed to have better keyword independence.

Here is an example using the older Draft 4 convention:

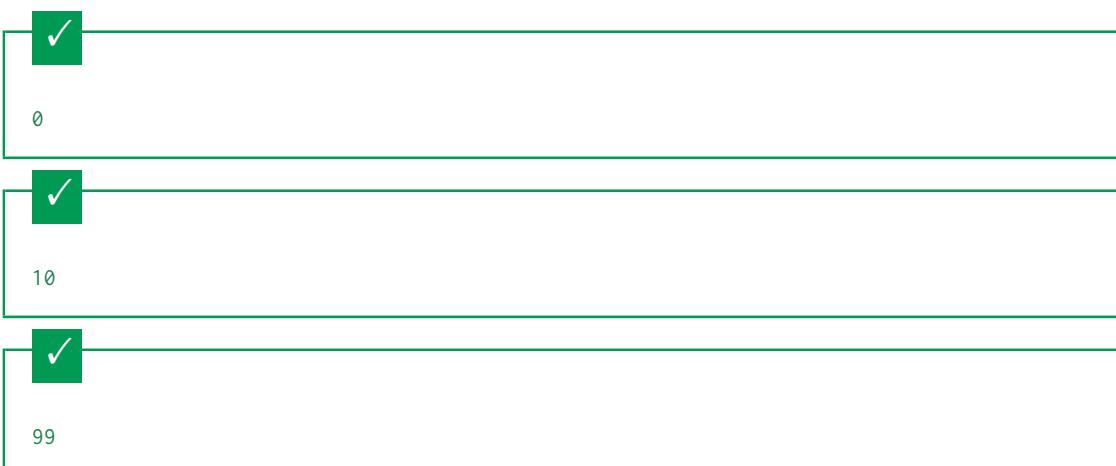
```
{ json schema }

{
  "type": "number",
  "minimum": 0,
  "maximum": 100,
  "exclusiveMaximum": true
}
```

Less than `minimum`:



`exclusiveMinimum` was not specified, so 0 is included:



`exclusiveMaximum` is `true`, so 100 is not included:



Greater than `maximum`:



4.5 object

- *Properties* (page 30)
- *Pattern Properties* (page 31)
- *Additional Properties* (page 32)
- *Required Properties* (page 34)
- *Property names* (page 36)
- *Size* (page 36)

Objects are the mapping type in JSON. They map “keys” to “values”. In JSON, the “keys” must always be strings. Each of these pairs is conventionally referred to as a “property”.

Python

In Python, "objects" are analogous to the `dict` type. An important difference, however, is that while Python dictionaries may use anything hashable as a key, in JSON all the keys must be strings.

Try not to be confused by the two uses of the word "object" here: Python uses the word `object` to mean the generic base class for everything, whereas in JSON it is used only to mean a mapping from string keys to values.

Ruby

In Ruby, "objects" are analogous to the `Hash` type. An important difference, however, is that all keys in JSON must be strings, and therefore any non-string keys are converted over to their string representation.

Try not to be confused by the two uses of the word "object" here: Ruby uses the word `Object` to mean the generic base class for everything, whereas in JSON it is used only to mean a mapping from string keys to values.

{ json schema }

```
{ "type": "object" }
```



```
{
  "key": "value",
  "another_key": "another_value"
}
```



```
{  
  "Sun": 1.9891e30,  
  "Jupiter": 1.8986e27,  
  "Saturn": 5.6846e26,  
  "Neptune": 10.243e25,  
  "Uranus": 8.6810e25,  
  "Earth": 5.9736e24,  
  "Venus": 4.8685e24,  
  "Mars": 6.4185e23,  
  "Mercury": 3.3022e23,  
  "Moon": 7.349e22,  
  "Pluto": 1.25e22  
}
```

Using non-strings as keys is invalid JSON:



```
{  
  0.01: "cm",  
  1: "m",  
  1000: "km"  
}
```



"Not an object"



["An", "array", "not", "an", "object"]

4.5.1 Properties

The properties (key-value pairs) on an object are defined using the `properties` keyword. The value of `properties` is an object, where each key is the name of a property and each value is a schema used to validate that property. Any property that doesn't match any of the property names in the `properties` keyword is ignored by this keyword.

Note: See [Additional Properties](#) (page 32) for how to disallow properties that don't match any of the property names in `properties`.

For example, let's say we want to define a simple schema for an address made up of a number, street name and street type:

{ json schema }

```
{
  "type": "object",
  "properties": {
    "number": { "type": "number" },
    "street_name": { "type": "string" },
    "street_type": { "enum": ["Street", "Avenue", "Boulevard"] }
  }
}
```



```
{ "number": 1600, "street_name": "Pennsylvania", "street_type": "Avenue" }
```

If we provide the number in the wrong type, it is invalid:



```
{ "number": "1600", "street_name": "Pennsylvania", "street_type": "Avenue" }
```

By default, leaving out properties is valid. See *Required Properties* (page 34).



```
{ "number": 1600, "street_name": "Pennsylvania" }
```

By extension, even an empty object is valid:



```
{ }
```

By default, providing additional properties is valid:



```
{ "number": 1600, "street_name": "Pennsylvania", "street_type": "Avenue", "direction": "NW" }
```

4.5.2 Pattern Properties

Sometimes you want to say that, given a particular kind of property name, the value should match a particular schema. That's where `patternProperties` comes in: it maps regular expressions to schemas. If a property name matches the given regular expression, the property value must validate against the corresponding schema.

Note: Regular expressions are not anchored. This means that when defining the regular expressions for `patternProperties`, it's important to note that the expression may match anywhere within the property name. For example, the regular expression "`p`" will match any property name with a `p` in it, such as "`apple`", not just a property whose name is simply "`p`". It's therefore usually less confusing to surround the regular expression in `^...$`,

for example, "`^p$`".

In this example, any properties whose names start with the prefix `S_` must be strings, and any with the prefix `I_` must be integers. Any properties that do not match either regular expression are ignored.

{ json schema }

```
{  
  "type": "object",  
  "patternProperties": {  
    "^S_": { "type": "string" },  
    "^I_": { "type": "integer" }  
  }  
}
```



```
{ "S_25": "This is a string" }
```



```
{ "I_0": 42 }
```

If the name starts with `S_`, it must be a string



```
{ "S_0": 42 }
```

If the name starts with `I_`, it must be an integer



```
{ "I_42": "This is a string" }
```

This is a key that doesn't match any of the regular expressions:



```
{ "keyword": "value" }
```

4.5.3 Additional Properties

The `additionalProperties` keyword is used to control the handling of extra stuff, that is, properties whose names are not listed in the `properties` keyword or match any of the regular expressions in the `patternProperties` keyword. By default any additional properties are allowed.

The value of the `additionalProperties` keyword is a schema that will be used to validate any properties in the instance that are not matched by `properties` or `patternProperties`. Setting the `additionalProperties` schema to

`false` means no additional properties will be allowed.

Reusing the example from *Properties* (page 30), but this time setting `additionalProperties` to `false`.

```
{json schema}
{
  "type": "object",
  "properties": {
    "number": { "type": "number" },
    "street_name": { "type": "string" },
    "street_type": { "enum": ["Street", "Avenue", "Boulevard"] }
  },
  "additionalProperties": false
}
```



```
{ "number": 1600, "street_name": "Pennsylvania", "street_type": "Avenue" }
```

Since `additionalProperties` is `false`, this extra property “`direction`” makes the object invalid:



```
{ "number": 1600, "street_name": "Pennsylvania", "street_type": "Avenue", "direction": "NW" }
```

You can use non-boolean schemas to put more complex constraints on the additional properties of an instance. For example, one can allow additional properties, but only if they are each a string:

```
{json schema}
{
  "type": "object",
  "properties": {
    "number": { "type": "number" },
    "street_name": { "type": "string" },
    "street_type": { "enum": ["Street", "Avenue", "Boulevard"] }
  },
  "additionalProperties": { "type": "string" }
}
```



```
{ "number": 1600, "street_name": "Pennsylvania", "street_type": "Avenue" }
```

This is valid, since the additional property’s value is a string:



```
{ "number": 1600, "street_name": "Pennsylvania", "street_type": "Avenue", "direction": "NW" }
```

This is invalid, since the additional property's value is not a string:



```
{ "number": 1600, "street_name": "Pennsylvania", "street_type": "Avenue", "office_number":  
→201 }
```

You can use `additionalProperties` with a combination of `properties` and `patternProperties`. In the following example, based on the example from [Pattern Properties](#) (page 31), we add a "builtin" property, which must be a number, and declare that all additional properties (that are neither defined by `properties` nor matched by `patternProperties`) must be strings:

{ json schema }

```
{  
  "type": "object",  
  "properties": {  
    "builtin": { "type": "number" }  
  },  
  "patternProperties": {  
    "^S_": { "type": "string" },  
    "^I_": { "type": "integer" }  
  },  
  "additionalProperties": { "type": "string" }  
}
```



```
{ "builtin": 42 }
```

This is a key that doesn't match any of the regular expressions:



```
{ "keyword": "value" }
```

It must be a string:



```
{ "keyword": 42 }
```

4.5.4 Required Properties

By default, the properties defined by the `properties` keyword are not required. However, one can provide a list of required properties using the `required` keyword.

The `required` keyword takes an array of zero or more strings. Each of these strings must be unique.

Draft 4

In Draft 4, required must contain at least one string.

In the following example schema defining a user record, we require that each user has a name and e-mail address, but we don't mind if they don't provide their address or telephone number:

{ json schema }

```
{
  "type": "object",
  "properties": {
    "name": { "type": "string" },
    "email": { "type": "string" },
    "address": { "type": "string" },
    "telephone": { "type": "string" }
  },
  "required": ["name", "email"]
}
```



```
{
  "name": "William Shakespeare",
  "email": "bill@stratford-upon-avon.co.uk"
}
```

Providing extra properties is fine, even properties not defined in the schema:



```
{
  "name": "William Shakespeare",
  "email": "bill@stratford-upon-avon.co.uk",
  "address": "Henley Street, Stratford-upon-Avon, Warwickshire, England",
  "authorship": "in question"
}
```

Missing the required "email" property makes the JSON document invalid:



```
{
  "name": "William Shakespeare",
  "address": "Henley Street, Stratford-upon-Avon, Warwickshire, England",
}
```

In JSON a property with value `null` is not equivalent to the property not being present. This fails because `null` is not of type "string", it's of type "null"



```
{  
  "name": "William Shakespeare",  
  "address": "Henley Street, Stratford-upon-Avon, Warwickshire, England",  
  "email": null  
}
```

4.5.5 Property names

New in draft 6

The names of properties can be validated against a schema, irrespective of their values. This can be useful if you don't want to enforce specific properties, but you want to make sure that the names of those properties follow a specific convention. You might, for example, want to enforce that all names are valid ASCII tokens so they can be used as attributes in a particular programming language.

{ json schema }

```
{  
  "type": "object",  
  "propertyNames": {  
    "pattern": "^[A-Za-z_][A-Za-z0-9_]*$"  
  }  
}
```



```
{  
  "_a_proper_token_001": "value"  
}
```



```
{  
  "001 invalid": "value"  
}
```

Since object keys must always be strings anyway, it is implied that the schema given to `propertyNames` is always at least:

```
{ "type": "string" }
```

4.5.6 Size

The number of properties on an object can be restricted using the `minProperties` and `maxProperties` keywords. Each of these must be a non-negative integer.

{ json schema }

```
{
  "type": "object",
  "minProperties": 2,
  "maxProperties": 3
}
```



{}



{ "a": 0 }



{ "a": 0, "b": 1 }



{ "a": 0, "b": 1, "c": 2 }



{ "a": 0, "b": 1, "c": 2, "d": 3 }

4.6 array

- *Items* (page 38)
 - *List validation* (page 38)
 - *Tuple validation* (page 39)
- *Additional Items* (page 40)
- *Contains* (page 42)
- *Length* (page 43)
- *Uniqueness* (page 44)

Arrays are used for ordered elements. In JSON, each element in an array may be of a different type.

Python

In Python, "array" is analogous to a `list` or `tuple` type, depending on usage. However, the `json` module in the Python standard library will always use Python lists to represent JSON arrays.

Ruby

In Ruby, "array" is analogous to a `Array` type.

{json schema}

```
{ "type": "array" }
```



[1, 2, 3, 4, 5]



[3, "different", { "types" : "of values" }]



{"Not": "an array"}

4.6.1 Items

There are two ways in which arrays are generally used in JSON:

- **List validation:** a sequence of arbitrary length where each item matches the same schema.
- **Tuple validation:** a sequence of fixed length where each item may have a different schema. In this usage, the index (or location) of each item is meaningful as to how the value is interpreted. (This usage is often given a whole separate type in some programming languages, such as Python's `tuple`).

List validation

List validation is useful for arrays of arbitrary length where each item matches the same schema. For this kind of array, set the `items` keyword to a single schema that will be used to validate all of the items in the array.

Note: When `items` is a single schema, the `additionalItems` keyword is meaningless, and it should not be used.

In the following example, we define that each item in an array is a number:

{ json schema }

```
{
  "type": "array",
  "items": {
    "type": "number"
  }
}
```



[1, 2, 3, 4, 5]

A single “non-number” causes the whole array to be invalid:



[1, 2, "3", 4, 5]

The empty array is always valid:



[]

Tuple validation

Tuple validation is useful when the array is a collection of items where each has a different schema and the ordinal index of each item is meaningful.

For example, you may represent a street address such as:

1600 Pennsylvania Avenue NW

as a 4-tuple of the form:

[number, street_name, street_type, direction]

Each of these fields will have a different schema:

- `number`: The address number. Must be a number.
- `street_name`: The name of the street. Must be a string.
- `street_type`: The type of street. Should be a string from a fixed set of values.
- `direction`: The city quadrant of the address. Should be a string from a different set of values.

To do this, we set the `items` keyword to an array, where each item is a schema that corresponds to each index of the document’s array. That is, an array where the first element validates the first element of the input array, the second element validates the second element of the input array, etc.

Here’s the example schema:

{ json schema }

```
{  
  "type": "array",  
  "items": [  
    { "type": "number" },  
    { "type": "string" },  
    { "enum": ["Street", "Avenue", "Boulevard"] },  
    { "enum": ["NW", "NE", "SW", "SE"] }  
  ]  
}
```



[1600, "Pennsylvania", "Avenue", "NW"]

“Drive” is not one of the acceptable street types:



[24, "Sussex", "Drive"]

This address is missing a street number



["Palais de l'Élysée"]

It's okay to not provide all of the items:



[10, "Downing", "Street"]

And, by default, it's also okay to add additional items to end:



[1600, "Pennsylvania", "Avenue", "NW", "Washington"]

4.6.2 Additional Items

The `additionalItems` keyword controls whether it's valid to have additional items in a tuple beyond what is defined in `items`. The value of the `additionalItems` keyword is a schema that all additional items must pass in order for the keyword to validate. This keyword is ignored if there is not a “tuple validation” `items` keyword present in the same schema.

Draft 4

In Draft 4, `additionalItems` does not require a "tuple validation" `items` keyword to be present. There are no constraints on any of the items, so all items are considered to be additional items.

Here, we'll reuse the example schema above, but set `additionalItems` to `false`, which has the effect of disallowing extra items in the array.

{ json schema }

```
{
  "type": "array",
  "items": [
    { "type": "number" },
    { "type": "string" },
    { "enum": ["Street", "Avenue", "Boulevard"] },
    { "enum": ["NW", "NE", "SW", "SE"] }
  ],
  "additionalItems": false
}
```



[1600, "Pennsylvania", "Avenue", "NW"]

It's ok to not provide all of the items:



[1600, "Pennsylvania", "Avenue"]

But, since `additionalItems` is `false`, we can't provide extra items:



[1600, "Pennsylvania", "Avenue", "NW", "Washington"]

You can express more complex constraints by using a non-boolean schema to constrain what value additional items can have. In that case, we could say that additional items are allowed, as long as they are all strings:

{ json schema }

```
{  
  "type": "array",  
  "items": [  
    { "type": "number" },  
    { "type": "string" },  
    { "enum": ["Street", "Avenue", "Boulevard"] },  
    { "enum": ["NW", "NE", "SW", "SE"] }  
,  
  "additionalItems": { "type": "string" }  
}
```

Extra string items are ok ...



```
[1600, "Pennsylvania", "Avenue", "NW", "Washington"]
```

... but not anything else



```
[1600, "Pennsylvania", "Avenue", "NW", 20500]
```

Note: Because “list validation” (`items` is an object) applies to all items in the list there are no additional items and therefore `additionalItems` has nothing to apply its schema to and will have no effect.

4.6.3 Contains

New in draft 6

While the `items` schema must be valid for every item in the array, the `contains` schema only needs to validate against one or more items in the array.

{ json schema }

```
{  
  "type": "array",  
  "contains": {  
    "type": "number"  
  }  
}
```

A single “number” is enough to make this pass:



```
["life", "universe", "everything", 42]
```

But if we have no number, it fails:



```
["life", "universe", "everything", "forty-two"]
```

All numbers is, of course, also okay:



```
[1, 2, 3, 4, 5]
```

4.6.4 Length

The length of the array can be specified using the `minItems` and `maxItems` keywords. The value of each keyword must be a non-negative number. These keywords work whether doing *List validation* (page 38) or *Tuple validation* (page 39).

{ json schema }

```
{
  "type": "array",
  "minItems": 2,
  "maxItems": 3
}
```



```
[]
```



```
[1]
```



```
[1, 2]
```



```
[1, 2, 3]
```



```
[1, 2, 3, 4]
```

4.6.5 Uniqueness

A schema can ensure that each of the items in an array is unique. Simply set the `uniqueItems` keyword to `true`.

{ json schema }

```
{
  "type": "array",
  "uniqueItems": true
}
```

✓

```
[1, 2, 3, 4, 5]
```

✗

```
[1, 2, 3, 3, 4]
```

The empty array always passes:

✓

```
[]
```

4.7 boolean

The `boolean` type matches only two special values: `true` and `false`. Note that values that *evaluate* to `true` or `false`, such as `1` and `0`, are not accepted by the schema.

Python

```
In Python, "boolean" is analogous to bool. Note that in JSON, true and false are lower case, whereas in Python they are capitalized (True and False).
```

Ruby

```
In Ruby, "boolean" is analogous to TrueClass and FalseClass. Note that in Ruby there is no Boolean class.
```

{ json schema }

```
{ "type": "boolean" }
```

	true
	false
	"true"

Values that evaluate to true or false are still not accepted by the schema:

	0
--	---

4.8 null

When a schema specifies a type of null, it has only one acceptable value: null.

Note: It's important to remember that in JSON, null isn't equivalent to something being absent. See [Required Properties](#) (page 34) for an example.

	Python
In Python, null is analogous to None.	

	Ruby
In Ruby, null is analogous to nil.	

	{ json schema }
{ "type": "null" }	
	null



4.9 Generic keywords

- *Annotations* (page 46)
- *Comments* (page 47)
- *Enumerated values* (page 47)
- *Constant values* (page 48)

This chapter lists some miscellaneous properties that are available for all JSON types.

4.9.1 Annotations

JSON Schema includes a few keywords, that aren't strictly used for validation, but are used to describe parts of a schema. None of these “annotation” keywords are required, but they are encouraged for good practice, and can make your schema “self-documenting”.

The `title` and `description` keywords must be strings. A “`title`” will preferably be short, whereas a “`description`” will provide a more lengthy explanation about the purpose of the data described by the schema.

The `default` keyword specifies a default value. This value is not used to fill in missing values during the validation process. Non-validation tools such as documentation generators or form generators may use this value to give hints to users about how to use a value. However, `default` is typically used to express that if a value is missing, then the value is semantically the same as if the value was present with the default value. The value of `default` should validate against the schema in which it resides, but that isn't required.

New in draft 6 The `examples` keyword is a place to provide an array of examples that validate against the schema. This isn't used for validation, but may help with explaining the effect and purpose of the schema to a reader. Each entry should validate against the schema in which it resides, but that isn't strictly required. There is no need to duplicate the `default` value in the `examples` array, since `default` will be treated as another example.

New in draft 7 The boolean keywords `readOnly` and `writeOnly` are typically used in an API context. `readOnly` indicates that a value should not be modified. It could be used to indicate that a PUT request that changes a value would result in a 400 Bad Request response. `writeOnly` indicates that a value may be set, but will remain hidden. It could be used to indicate you can set a value with a PUT request, but it would not be included when retrieving that record with a GET request.

```
{  
  "title": "Match anything",  
  "description": "This is a schema that matches anything.",  
  "default": "Default value",  
  "examples": [  
    "Anything",  
    4035  
,  
    "readOnly": true,  
    "writeOnly": false  
}
```

4.9.2 Comments

New in draft 7 `$comment`

The `$comment` keyword is strictly intended for adding comments to a schema. Its value must always be a string. Unlike the annotations `title`, `description`, and `examples`, JSON schema implementations aren't allowed to attach any meaning or behavior to it whatsoever, and may even strip them at any time. Therefore, they are useful for leaving notes to future editors of a JSON schema, but should not be used to communicate to users of the schema.

4.9.3 Enumerated values

The `enum` keyword is used to restrict a value to a fixed set of values. It must be an array with at least one element, where each element is unique.

The following is an example for validating street light colors:

```
{  
  "enum": ["red", "amber", "green"]  
}
```



"red"



"blue"

You can use `enum` even without a type, to accept values of different types. Let's extend the example to use `null` to indicate "off", and also add `42`, just for fun.

{ json schema }

```
{ "enum": [ "red", "amber", "green", null, 42 ] }
```


"red"


null


42


0

4.9.4 Constant values

New in draft 6

The `const` keyword is used to restrict a value to a single value.

For example, if you only support shipping to the United States for export reasons:

{ json schema }

```
{ "properties": { "country": { "const": "United States of America" } } }
```


{ "country": "United States of America" }



```
{ "country": "Canada" }
```

4.10 Media: string-encoding non-JSON data

- [contentMediaType](#) (page 49)
- [contentEncoding](#) (page 49)
- [Examples](#) (page 49)

New in draft 7

JSON schema has a set of keywords to describe and optionally validate non-JSON data stored inside JSON strings. Since it would be difficult to write validators for many media types, JSON schema validators are not required to validate the contents of JSON strings based on these keywords. However, these keywords are still useful for an application that consumes validated JSON.

4.10.1 contentMediaType

The `contentMediaType` keyword specifies the MIME type of the contents of a string, as described in [RFC 2046](#). There is a list of [MIME types officially registered by the IANA](#), but the set of types supported will be application and operating system dependent. Mozilla Developer Network also maintains a [shorter list of MIME types that are important for the web](#).

4.10.2 contentEncoding

The `contentEncoding` keyword specifies the encoding used to store the contents, as specified in [RFC 2054](#), part 6.1. The acceptable values are `7bit`, `8bit`, `binary`, `quoted-printable` and `base64`. If not specified, the encoding is the same as the containing JSON document.

Without getting into the low-level details of each of these encodings, there are really only two options useful for modern usage:

- If the content is encoded in the same encoding as the enclosing JSON document (which for practical purposes, is almost always UTF-8), leave `contentEncoding` unspecified, and include the content in a string as-is. This includes text-based content types, such as `text/html` or `application/xml`.
- If the content is binary data, set `contentEncoding` to `base64` and encode the contents using [Base64](#). This would include many image types, such as `image/png` or audio types, such as `audio/mpeg`.

4.10.3 Examples

The following schema indicates the string contains an HTML document, encoded using the same encoding as the surrounding document:

{ json schema }

```
{  
  "type": "string",  
  "contentMediaType": "text/html"  
}
```



```
"<!DOCTYPE html><html xmlns=\\"http://www.w3.org/1999/xhtml\\"><head></head></html>"
```

The following schema indicates that a string contains a **PNG** image, encoded using Base64:

{ json schema }

```
{  
  "type": "string",  
  "contentEncoding": "base64",  
  "contentMediaType": "image/png"  
}
```



```
"iVBORw0KGgoAAAANSUhEUgAAABgAAAAACAYAAADgdz34AAAABmJLR0QA/wD/AP+gvaeTAAAA..."
```

4.11 Schema Composition

- *allOf* (page 51)
- *anyOf* (page 51)
- *oneOf* (page 52)
- *not* (page 53)
- *Properties of Schema Composition* (page 53)
 - *Subschema Independence* (page 53)
 - *Illogical Schemas* (page 55)
 - *Factoring Schemas* (page 56)

JSON Schema includes a few keywords for combining schemas together. Note that this doesn't necessarily mean combining schemas from multiple files or JSON trees, though these facilities help to enable that and are described in *Structuring a complex schema* (page 71). Combining schemas may be as simple as allowing a value to be validated against multiple criteria at the same time.

These keywords correspond to well known boolean algebra concepts like AND, OR, XOR, and NOT. You can often use these keywords to express complex constraints that can't otherwise be expressed with standard JSON Schema keywords.

The keywords used to combine schemas are:

- *allOf* (page 51): (AND) Must be valid against *all* of the subschemas
- *anyOf* (page 51): (OR) Must be valid against *any* of the subschemas
- *oneOf* (page 52): (XOR) Must be valid against *exactly one* of the subschemas

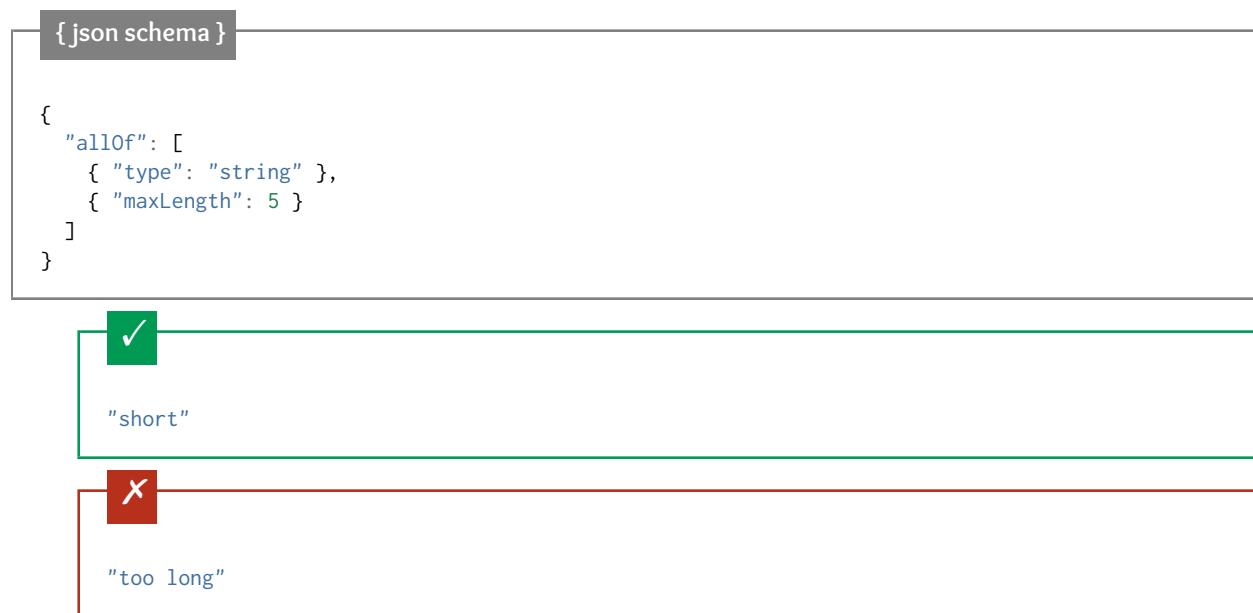
All of these keywords must be set to an array, where each item is a schema.

In addition, there is:

- *not* (page 53): (NOT) Must *not* be valid against the given schema

4.11.1 allOf

To validate against `allOf`, the given data must be valid against all of the given subschemas.



Note: *allOf* (page 51) can not be used to “extend” a schema to add more details to it in the sense of object-oriented inheritance. Instances must independently be valid against “all of” the schemas in the `allOf`. See the section on *Subschema Independence* (page 53) for more information.

4.11.2 anyOf

To validate against `anyOf`, the given data must be valid against any (one or more) of the given subschemas.

{ json schema }

```
{  
  "anyOf": [  
    { "type": "string", "maxLength": 5 },  
    { "type": "number", "minimum": 0 }  
  ]  
}
```



"short"



"too long"



12



-5

4.11.3 oneOf

To validate against oneOf, the given data must be valid against exactly one of the given subschemas.

{ json schema }

```
{  
  "oneOf": [  
    { "type": "number", "multipleOf": 5 },  
    { "type": "number", "multipleOf": 3 }  
  ]  
}
```



10



9

Not a multiple of either 5 or 3.



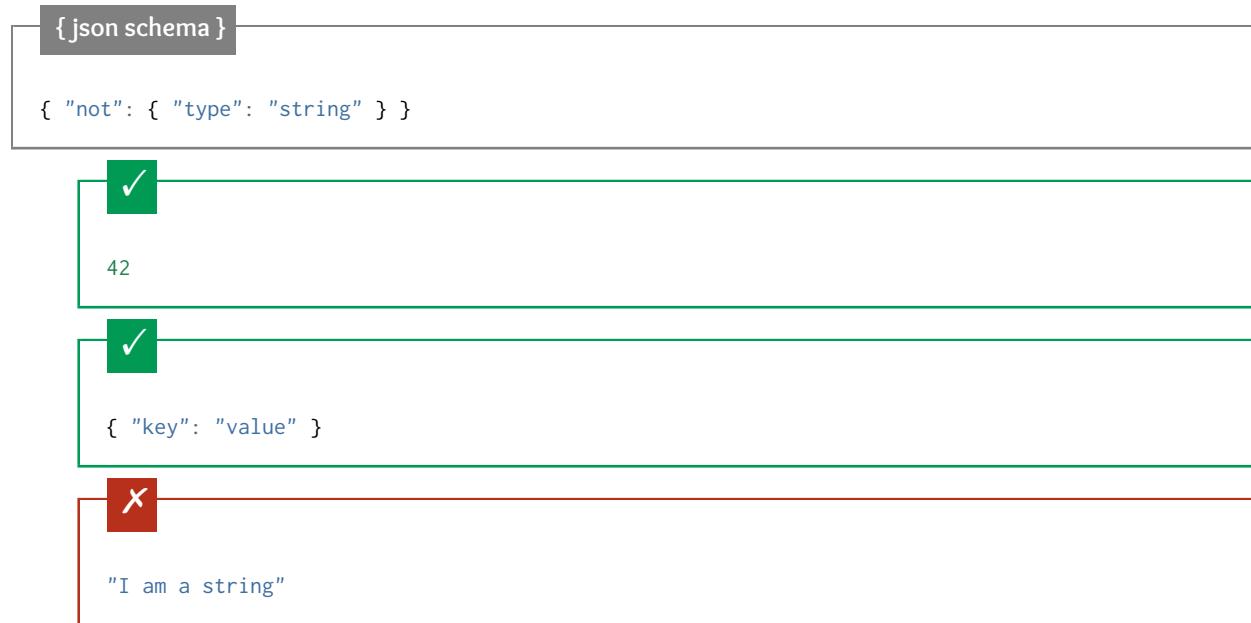
Multiple of *both* 5 and 3 is rejected.



4.11.4 not

The `not` keyword declares that an instance validates if it doesn't validate against the given subschema.

For example, the following schema validates against anything that is not a string:



4.11.5 Properties of Schema Composition

Subschema Independence

It is important to note that the schemas listed in an `allOf` (page 51), `anyOf` (page 51) or `oneOf` (page 52) array know nothing of one another. For example, say you had a schema for an address in a definitions section, and want to “extend” it to include an address type:

{ json schema }

```
{  
  "definitions": {  
    "address": {  
      "type": "object",  
      "properties": {  
        "street_address": { "type": "string" },  
        "city": { "type": "string" },  
        "state": { "type": "string" }  
      },  
      "required": ["street_address", "city", "state"]  
    }  
  },  
  "allOf": [  
    { "$ref": "#/definitions/address" },  
    {  
      "properties": {  
        "type": { "enum": [ "residential", "business" ] }  
      }  
    }  
  ]  
}
```



```
{  
  "street_address": "1600 Pennsylvania Avenue NW",  
  "city": "Washington",  
  "state": "DC",  
  "type": "business"  
}
```

This works, but what if we wanted to restrict the schema so no additional properties are allowed? One might try adding the highlighted line below:

{ json schema }

```
{
  "definitions": {
    "address": {
      "type": "object",
      "properties": {
        "street_address": { "type": "string" },
        "city": { "type": "string" },
        "state": { "type": "string" }
      },
      "required": ["street_address", "city", "state"]
    }
  },
  "allOf": [
    { "$ref": "#/definitions/address" },
    {
      "properties": {
        "type": { "enum": [ "residential", "business" ] }
      }
    }
  ],
  "additionalProperties": false
}
```



```
{
  "street_address": "1600 Pennsylvania Avenue NW",
  "city": "Washington",
  "state": "DC",
  "type": "business"
}
```

Unfortunately, now the schema will reject *everything*. This is because `additionalProperties` knows nothing about the properties declared in the subschemas inside of the `allOf` (page 51) array.

To many, this is one of the biggest surprises of the combining operations in JSON schema: it does not behave like inheritance in an object-oriented language. There are some proposals to address this in the next version of the JSON schema specification.

Illogical Schemas

Note that it's quite easy to create schemas that are logical impossibilities with these keywords. The following example creates a schema that won't validate against anything (since something may not be both a string and a number at the same time):

{ json schema }

```
{  
  "allOf": [  
    { "type": "string" },  
    { "type": "number" }  
  ]  
}
```

X

"No way"

X

-1

Factoring Schemas

Note that it's possible to "factor" out the common parts of the subschemas. The following two schemas are equivalent.

{ json schema }

```
{  
  "oneOf": [  
    { "type": "number", "multipleOf": 5 },  
    { "type": "number", "multipleOf": 3 }  
  ]  
}
```

{ json schema }

```
{  
  "type": "number",  
  "oneOf": [  
    { "multipleOf": 5 },  
    { "multipleOf": 3 }  
  ]  
}
```

4.12 Applying subschemas conditionally

- *Dependencies* (page 57)
 - *Property dependencies* (page 57)
 - *Schema dependencies* (page 59)
- *If-Then-Else* (page 61)
- *Implication* (page 65)

4.12.1 Dependencies

The `dependencies` keyword conditionally applies additional constraints to a schema based on the presence of certain properties. There are two forms of dependencies in JSON Schema:

- **Property dependencies** declare that certain other properties must be present if a given property is present.
- **Schema dependencies** declare that the schema changes when a given property is present.

Property dependencies

Let's start with the simpler case of property dependencies. For example, suppose we have a schema representing a customer. If you have their credit card number, you also want to ensure you have a billing address. If you don't have their credit card number, a billing address would not be required. We represent this dependency of one property on another using the `dependencies` keyword. The value of the `dependencies` keyword is an object. Each entry in the object maps from the name of a property, p , to an array of strings listing properties that are required if p is present.

In the following example, whenever a `credit_card` property is provided, a `billing_address` property must also be present:

```
{ json schema }

{
  "type": "object",

  "properties": {
    "name": { "type": "string" },
    "credit_card": { "type": "number" },
    "billing_address": { "type": "string" }
  },
  "required": ["name"],

  "dependencies": {
    "credit_card": ["billing_address"]
  }
}
```



```
{  
  "name": "John Doe",  
  "credit_card": 5555555555555555,  
  "billing_address": "555 Debtor's Lane"  
}
```

This instance has a `credit_card`, but it's missing a `billing_address`.



```
{  
  "name": "John Doe",  
  "credit_card": 5555555555555555  
}
```

This is okay, since we have neither a `credit_card`, or a `billing_address`.



```
{  
  "name": "John Doe"  
}
```

Note that dependencies are not bidirectional. It's okay to have a billing address without a credit card number.



```
{  
  "name": "John Doe",  
  "billing_address": "555 Debtor's Lane"  
}
```

To fix the last issue above (that dependencies are not bidirectional), you can, of course, define the bidirectional dependencies explicitly:

{ json schema }

```
{
  "type": "object",

  "properties": {
    "name": { "type": "string" },
    "credit_card": { "type": "number" },
    "billing_address": { "type": "string" }
  },

  "required": ["name"],

  "dependencies": {
    "credit_card": ["billing_address"],
    "billing_address": ["credit_card"]
  }
}
```

This instance has a `credit_card`, but it's missing a `billing_address`.



```
{
  "name": "John Doe",
  "credit_card": 5555555555555555
}
```

This has a `billing_address`, but is missing a `credit_card`.



```
{
  "name": "John Doe",
  "billing_address": "555 Debtor's Lane"
}
```

Schema dependencies

Schema dependencies work like property dependencies, but instead of just specifying additional required properties, it specifies a schema that will be applied to the instance. This schema is applied in the same way [allOf](#) (page 51) applies schemas. Nothing is merged or extended. Both schemas apply independently.

For example, here is another way to write the above:

{ json schema }

```
{
  "type": "object",

  "properties": {
    "name": { "type": "string" },
    "credit_card": { "type": "number" }
  },

  "required": ["name"],

  "dependencies": {
    "credit_card": {
      "properties": {
        "billing_address": { "type": "string" }
      },
      "required": ["billing_address"]
    }
  }
}
```



```
{
  "name": "John Doe",
  "credit_card": 5555555555555555,
  "billing_address": "555 Debtor's Lane"
}
```

This instance has a `credit_card`, but it's missing a `billing_address`:



```
{
  "name": "John Doe",
  "credit_card": 5555555555555555
}
```

This has a `billing_address`, but is missing a `credit_card`. This passes, because here `billing_address` just looks like an additional property:



```
{
  "name": "John Doe",
  "billing_address": "555 Debtor's Lane"
}
```

4.12.2 If-Then-Else

New in draft 7 The `if`, `then` and `else` keywords allow the application of a subschema based on the outcome of another schema, much like the `if/then/else` constructs you've probably seen in traditional programming languages.

If `if` is valid, `then` must also be valid (and `else` is ignored.) If `if` is invalid, `else` must also be valid (and `then` is ignored).

If `then` or `else` is not defined, `if` behaves as if they have a value of `true`.

If `then` and/or `else` appear in a schema without `if`, `then` and `else` are ignored.

We can put this in the form of a truth table, showing the combinations of when `if`, `then`, and `else` are valid and the resulting validity of the entire schema:

if	then	else	whole schema
T	T	n/a	T
T	F	n/a	F
F	n/a	T	T
F	n/a	F	F
n/a	n/a	n/a	T

For example, let's say you wanted to write a schema to handle addresses in the United States and Canada. These countries have different postal code formats, and we want to select which format to validate against based on the country. If the address is in the United States, the `postal_code` field is a "zipcode": five numeric digits followed by an optional four digit suffix. If the address is in Canada, the `postal_code` field is a six digit alphanumeric string where letters and numbers alternate.

```
{ json schema }
{
  "type": "object",
  "properties": {
    "street_address": {
      "type": "string"
    },
    "country": {
      "default": "United States of America",
      "enum": ["United States of America", "Canada"]
    }
  },
  "if": {
    "properties": { "country": { "const": "United States of America" } }
  },
  "then": {
    "properties": { "postal_code": { "pattern": "[0-9]{5}(-[0-9]{4})?" } }
  },
  "else": {
    "properties": { "postal_code": { "pattern": "[A-Z][0-9][A-Z] [0-9][A-Z][0-9]" } }
  }
}
```



```
{  
  "street_address": "1600 Pennsylvania Avenue NW",  
  "country": "United States of America",  
  "postal_code": "20500"  
}
```



```
{  
  "street_address": "1600 Pennsylvania Avenue NW",  
  "postal_code": "20500"  
}
```



```
{  
  "street_address": "24 Sussex Drive",  
  "country": "Canada",  
  "postal_code": "K1M 1M4"  
}
```



```
{  
  "street_address": "24 Sussex Drive",  
  "country": "Canada",  
  "postal_code": "10000"  
}
```



```
{  
  "street_address": "1600 Pennsylvania Avenue NW",  
  "postal_code": "K1M 1M4"  
}
```

Note: In this example, “country” is not a required property. Because the “if” schema also doesn’t require the “country” property, it will pass and the “then” schema will apply. Therefore, if the “country” property is not defined, the default behavior is to validate “postal_code” as a USA postal code. The “default” keyword doesn’t have an effect, but is nice to include for readers of the schema to more easily recognize the default behavior.

Unfortunately, this approach above doesn’t scale to more than two countries. You can, however, wrap pairs of if and then inside an allOf to create something that would scale. In this example, we’ll use United States and Canadian postal codes, but also add Netherlands postal codes, which are 4 digits followed by two letters. It’s left as an exercise to the reader to expand this to the remaining postal codes of the world.

{ json schema }

```
{
  "type": "object",
  "properties": {
    "street_address": {
      "type": "string"
    },
    "country": {
      "default": "United States of America",
      "enum": ["United States of America", "Canada", "Netherlands"]
    }
  },
  "allOf": [
    {
      "if": {
        "properties": { "country": { "const": "United States of America" } }
      },
      "then": {
        "properties": { "postal_code": { "pattern": "[0-9]{5}(-[0-9]{4})?" } }
      }
    },
    {
      "if": {
        "properties": { "country": { "const": "Canada" } },
        "required": ["country"]
      },
      "then": {
        "properties": { "postal_code": { "pattern": "[A-Z][0-9][A-Z] [0-9][A-Z][0-9]" } }
      }
    },
    {
      "if": {
        "properties": { "country": { "const": "Netherlands" } },
        "required": ["country"]
      },
      "then": {
        "properties": { "postal_code": { "pattern": "[0-9]{4} [A-Z]{2}" } }
      }
    }
  ]
}
```



```
{
  "street_address": "1600 Pennsylvania Avenue NW",
  "country": "United States of America",
  "postal_code": "20500"
}
```



```
{  
  "street_address": "1600 Pennsylvania Avenue NW",  
  "postal_code": "20500"  
}
```



```
{  
  "street_address": "24 Sussex Drive",  
  "country": "Canada",  
  "postal_code": "K1M 1M4"  
}
```



```
{  
  "street_address": "Adriaan Goekooplaan",  
  "country": "Netherlands",  
  "postal_code": "2517 JX"  
}
```



```
{  
  "street_address": "24 Sussex Drive",  
  "country": "Canada",  
  "postal_code": "10000"  
}
```



```
{  
  "street_address": "1600 Pennsylvania Avenue NW",  
  "postal_code": "K1M 1M4"  
}
```

Note: The “required” keyword is necessary in the “if” schemas or they would all apply if the “country” is not defined. Leaving “required” off of the “United States of America” “if” schema makes it effectively the default if no “country” is defined.

Note: Even if “country” was a required field, it’s still recommended to have the “required” keyword in each “if” schema. The validation result will be the same because “required” will fail, but not including it will add noise to error results because it will validate the “postal_code” against all three of the “then” schemas leading to irrelevant errors.

4.12.3 Implication

Before Draft 7, you can express an “if-then” conditional using the [Schema Composition](#) (page 50) keywords and a boolean algebra concept called “implication”. $A \rightarrow B$ (pronounced, A implies B) means that if A is true, then B must also be true. It can be expressed as $\neg A \mid\mid B$ which can be expressed as a JSON Schema.

{ json schema }

```
{
  "type": "object",
  "properties": {
    "restaurauntType": { "enum": ["fast-food", "sit-down"] },
    "total": { "type": "number" },
    "tip": { "type": "number" }
  },
  "anyOf": [
    {
      "not": {
        "properties": { "restaurantType": { "const": "sit-down" } },
        "required": ["restaurantType"]
      }
    },
    { "required": ["tip"] }
  ]
}
```

✓

```
{
  "restaurantType": "sit-down",
  "total": 16.99,
  "tip": 3.4
}
```

✗

```
{
  "restaurantType": "sit-down",
  "total": 16.99
}
```

✓

```
{
  "restaurantType": "fast-food",
  "total": 6.99
}
```

✓

```
{ "total": 5.25 }
```

Variations of implication can be used to express the same things you can express with the if/then/else keywords. if/then can be expressed as A → B, if/else can be expressed as !A → B, and if/then/else can be expressed as A → B AND !A → C.

Note: Since this pattern is not very intuitive, it's recommended to put your conditionals in definitions with a descriptive name and \$ref it into your schema with "allOf": [{ "\$ref": "#/definitions/sit-down-restaurant-implies-tip-is-required" }].

4.13 Declaring a Specification Version

- [\\$schema](#) (page 66)
- [Extending JSON Schema](#) (page 67)

Not only have there have been multiple versions of the specification, but it's also possible to add your own extensions to create your own version. JSON Schema provides ways for you to declare which version a schema conforms to and provides ways to describe your custom keywords.

4.13.1 \$schema

The \$schema keyword is used to declare which version of the JSON Schema specification the schema was written for. The value of the \$schema keyword is also the identifier for a schema that can be used to verify that the schema is valid according to the specification version \$schema identifies. A schema that describes another schema is called a “meta-schema”.

\$schema applies to the entire document and must be at the root level. It does not apply to externally referenced (\$ref) documents. Those schemas need to declare their own \$schema.

If \$schema is not used, an implementation might allow you to specify a value externally or it might make assumptions about which specification version should be used to evaluate the schema. It's recommended that all JSON Schemas have a \$schema entry to communicate to readers and tooling which specification version is intended. Therefore most of the time, you'll want this at the root of your schema:

```
"$schema": "http://json-schema.org/draft-07/schema#"
```

Draft-4

The identifier for Draft 4 is <http://json-schema.org/draft-04/schema#>.

Draft 4 defined a value for \$schema without a specific version (<http://json-schema.org/schema#>) which meant, use the latest version. This has since been deprecated and should no longer be used.

You might come across references to Draft 5. There is no Draft 5 release of JSON Schema. Draft 5 refers to a no-change revision of the Draft 4 release. It does not add, remove, or change any functionality. It only updates references, makes clarifications, and fixes bugs. Draft 5 describes the Draft 4 release. If you came here looking for information about Draft 5, you'll find it under Draft 4. We no longer use the "draft" terminology to refer to patch releases to avoid this confusion.

Draft-6

The identifier for Draft 6 is <http://json-schema.org/draft-06/schema#>.

4.13.2 Extending JSON Schema

You can extend the JSON Schema language to include your own custom keywords. To do this, you need to create a meta-schema that includes your custom keywords. The best way to do this is to make a copy of the meta-schema for the version you want to extend and make your changes to your copy. You will need to choose a custom URI to identify your custom version. This URI must not be one of the URIs used to identify official JSON Schema specification drafts and should probably include a domain name you own. You can use this URI with the `$schema` keyword to declare that your schemas use your custom version.

Note: Not all implementations support extending JSON Schema.

One of the strengths of JSON Schema is that it can be written in JSON and used in a variety of environments. For example, it can be used for both front-end and back-end HTML Form validation. The problem with adding custom keywords is that every environment where you want to use your schemas needs to understand how to evaluate your custom keywords. Meta-schemas can be used to ensure that schemas are written correctly, but each implementation will need custom code to understand how to evaluate the custom keywords.

Meta-data keywords are the most interoperable because they don't affect validation. For example, you could add a `units` keyword. This will always work as expected with a compliant validator.

{ json schema }

```
{
  "type": "number",
  "units": "kg"
}
```

✓

42

✗

"42"

The next best candidates for custom keywords are keywords that don't apply other schemas and don't modify the behavior of existing keywords. An `isEven` keyword is an example. In contexts where some validation is better than no validation such as validating an HTML Form in the browser, this schema will perform as well as can be expected. Full validation would still be required and should use a validator that understands the custom keyword.

{ json schema }

```
{  
  "type": "integer",  
  "isEven": true  
}
```



2

This passes because the validator doesn't understand isEven



3

The schema isn't completely impaired because it doesn't understand isEven



"3"

The least interoperable type of custom keyword is one that applies other schemas or modifies the behavior of existing keywords. An example would be something like `requiredProperties` that declares properties and makes them required. This example shows how the schema becomes almost completely useless when evaluated with a validator that doesn't understand the custom keyword. That doesn't necessarily mean that `requiredProperties` is a bad idea for a keyword, it's just not the right choice if the schema might need to be used in a context that doesn't understand custom keywords.

{ json schema }

```
{  
  "type": "object",  
  "requiredProperties": {  
    "foo": { "type": "string" }  
  }  
}
```



{ "foo": "bar" }

This passes because `requiredProperties` is not understood



{}

This passes because `requiredProperties` is not understood



Structuring a complex schema

- *Schema Identification* (page 71)
 - *JSON Pointer* (page 72)
 - *Named Anchors* (page 72)
- *Base URI* (page 73)
 - *Retrieval URI* (page 73)
 - *\$id* (page 74)
- *\$ref* (page 75)
- *definitions* (page 76)
- *Recursion* (page 77)
- *Bundling* (page 78)

When writing computer programs of even moderate complexity, it's commonly accepted that "structuring" the program into reusable functions is better than copying-and-pasting duplicate bits of code everywhere they are used. Likewise in JSON Schema, for anything but the most trivial schema, it's really useful to structure the schema into parts that can be reused in a number of places. This chapter will present the tools available for reusing and structuring schemas as well as some practical examples that use those tools.

5.1 Schema Identification

Like any other code, schemas are easier to maintain if they can be broken down into logical units that reference each other as necessary. In order to reference a schema, we need a way to identify a schema. Schema documents are identified by non-relative URIs.

Schema documents are not required to have an identifier, but you will need one if you want to reference one schema from another. In this document, we will refer to schemas with no identifier as "anonymous schemas".

In the following sections we will see how the “identifier” for a schema is determined.

Note: URI terminology can sometimes be unintuitive. In this document, the following definitions are used.

- **URI [1]** or **non-relative URI**: A full URI containing a scheme (`https`). It may contain a URI fragment (#foo). Sometimes this document will use “non-relative URI” to make it extra clear that relative URIs are not allowed.
 - **relative reference [2]**: A partial URI that does not contain a scheme (`https`). It may contain a fragment (#foo).
 - **URI-reference [3]**: A relative reference or non-relative URI. It may contain a URI fragment (#foo).
 - **absolute URI [4]** A full URI containing a scheme (`https`) but not a URI fragment (#foo).
-

Note: Even though schemas are identified by URIs, those identifiers are not necessarily network-addressable. They are just identifiers. Generally, implementations don’t make HTTP requests (`https://`) or read from the file system (`file://`) to fetch schemas. Instead, they provide a way to load schemas into an internal schema database. When a schema is referenced by its URI identifier, the schema is retrieved from the internal schema database.

5.1.1 JSON Pointer

In addition to identifying a schema document, you can also identify subschemas. The most common way to do that is to use a **JSON Pointer** in the URI fragment that points to the subschema.

A JSON Pointer describes a slash-separated path to traverse the keys in the objects in the document. Therefore, /`properties/street_address` means:

- 1) find the value of the key `properties`
- 2) within that object, find the value of the key `street_address`

The URI `https://example.com/schemas/address#/properties/street_address` identifies the highlighted subschema in the following schema.

{ json schema }

```
{
  "$id": "https://example.com/schemas/address",
  "type": "object",
  "properties": {
    "street_address": {
      "type": "string"
    },
    "city": {
      "type": "string"
    },
    "state": {
      "type": "string"
    }
  },
  "required": [
    "street_address",
    "city",
    "state"
  ]
}
```

5.1.2 Named Anchors

A less common way to identify a subschema is to create a named anchor in the schema using the `$id` keyword and using that name in the URI fragment. When the `$id` keyword contains a URI fragment, the fragment defines a named

anchor using the value of the fragment. Named anchors must start with a letter followed by any number of letters, digits, -, _, :, or ..

Draft 4

In Draft 4, `$id` is just `id` (without the dollar sign).

Note: If a named anchor is defined that doesn't follow these naming rules, then behavior is undefined. Your anchors might work in some implementation, but not others.

The URI `https://example.com/schemas/address#street_address` identifies the subschema on the highlighted part of the following schema.

{ json schema }

```
{
  "$id": "https://example.com/schemas/address",
  "type": "object",
  "properties": {
    "street_address": {
      "$id": "#street_address",
      "type": "string"
    },
    "city": { "type": "string" },
    "state": { "type": "string" }
  },
  "required": ["street_address", "city", "state"]
}
```

Note: JSON Schema doesn't define how `$id` should be interpreted when it contains both fragment and non-fragment URI parts. Therefore, when setting a named anchor, you should not use non-fragment URI parts in the URI-reference.

5.2 Base URI

Using non-relative URIs can be cumbersome, so any URIs used in JSON Schema can be URI-references that resolve against the schema's base URI resulting in a non-relative URI. This section describes how a schema's base URI is determined.

Note: Base URI determination and relative reference resolution is defined by RFC-3986. If you are familiar with how this works in HTML, this section should feel very familiar.

5.2.1 Retrieval URI

The URI used to fetch a schema is known as the “retrieval URI”. It's often possible to pass an anonymous schema to an implementation in which case that schema would have no retrieval URI.

Let's assume a schema is referenced using the URI <https://example.com/schemas/address> and the following schema is retrieved.

```
{ json schema }  
  
{  
  "type": "object",  
  "properties": {  
    "street_address": { "type": "string" },  
    "city": { "type": "string" },  
    "state": { "type": "string" }  
  },  
  "required": ["street_address", "city", "state"]  
}
```

The base URI for this schema is the same as the retrieval URI, <https://example.com/schemas/address>.

5.2.2 \$id

You can set the base URI using the `$id` keyword. The value of `$id` is a URI-reference that resolves against the *Retrieval URI* (page 73). The resulting URI is the base URI for the schema.

Draft 4

In Draft 4, `$id` is just `id` (without the dollar sign).

Note: This is analogous to the `<base>` tag [HTML](#).

Let's assume the URI <https://example.com/schema/address> and <https://example.com/schema/billing-address> both identify the following schema.

```
{ json schema }  
  
{  
  "$id": "/schemas/address",  
  
  "type": "object",  
  "properties": {  
    "street_address": { "type": "string" },  
    "city": { "type": "string" },  
    "state": { "type": "string" }  
  },  
  "required": ["street_address", "city", "state"]  
}
```

No matter which of the two URIs is used to retrieve this schema, the base URI will be <https://example.com/schemas/address>, which is the result of the `$id` URI-reference resolving against the *Retrieval URI* (page 73).

However, using a relative reference when setting a base URI can be problematic. For example, we couldn't use this schema as an anonymous schema because there would be no *Retrieval URI* (page 73) and you can't resolve a relative reference against nothing. For this and other reasons, it's recommended that you always use an absolute URI when declaring a base URI with `$id`.

The base URI of the following schema will always be <https://example.com/schemas/address> no matter what the *Retrieval URI* (page 73) was or if it's used as an anonymous schema.

```
{json schema}
{
  "$id": "https://example.com/schemas/address",
  "type": "object",
  "properties": {
    "street_address": { "type": "string" },
    "city": { "type": "string" },
    "state": { "type": "string" }
  },
  "required": ["street_address", "city", "state"]
}
```

Note: The behavior when setting a base URI that contains a URI fragment is undefined and should not be used because implementations may treat them differently.

5.3 \$ref

A schema can reference another schema using the `$ref` keyword. The value of `$ref` is a URI-reference that is resolved against the schema's *Base URI* (page 73). When evaluating a schema, an implementation uses the resolved identifier to retrieve the referenced schema and evaluation is continued from the retrieved schema.

`$ref` can be used anywhere a schema is expected. When an object contains a `$ref` property, the object is considered a reference, not a schema. Therefore, any other properties you put there will not be treated as JSON Schema keywords and will be ignored by the validator.

For this example, let's say we want to define a customer record, where each customer may have both a shipping and a billing address. Addresses are always the same—they have a street address, city and state—so we don't want to duplicate that part of the schema everywhere we want to store an address. Not only would that make the schema more verbose, but it makes updating it in the future more difficult. If our imaginary company were to start doing international business in the future and we wanted to add a country field to all the addresses, it would be better to do this in a single place rather than everywhere that addresses are used.

```
{json schema}
{
  "$id": "https://example.com/schemas/customer",
  "type": "object",
  "properties": {
    "first_name": { "type": "string" },
    "last_name": { "type": "string" },
    "shipping_address": { "$ref": "/schemas/address" },
    "billing_address": { "$ref": "/schemas/address" }
  },
  "required": ["first_name", "last_name", "shipping_address", "billing_address"]
}
```

The URI-references in \$ref resolve against the schema's *Base URI* (page 73) (<https://example.com/schemas/customer>) which results in <https://example.com/schemas/address>. The implementation retrieves that schema and uses it to evaluate the "shipping_address" and "billing_address" properties.

Note: When using \$ref in an anonymous schema, relative references may not be resolvable. Let's assume this example is used as an anonymous schema.

{ json schema }

```
{  
  "type": "object",  
  "properties": {  
    "first_name": { "type": "string" },  
    "last_name": { "type": "string" },  
    "shipping_address": { "$ref": "https://example.com/schemas/address" },  
    "billing_address": { "$ref": "/schemas/address" }  
  },  
  "required": ["first_name", "last_name", "shipping_address", "billing_address"]  
}
```

The \$ref at /properties/shipping_address can resolve just fine without a non-relative base URI to resolve against, but the \$ref at /properties/billing_address can't resolve to a non-relative URI and therefore can't be used to retrieve the address schema.

5.4 definitions

Sometimes we have small subschemas that are only intended for use in the current schema and it doesn't make sense to define them as separate schemas. Although we can identify any subschema using JSON Pointers or named anchors, the `definitions` keyword gives us a standardized place to keep subschemas intended for reuse in the current schema document.

Let's extend the previous customer schema example to use a common schema for the name properties. It doesn't make sense to define a new schema for this and it will only be used in this schema, so it's a good candidate for using `definitions`.

{ json schema }

```
{
  "$id": "https://example.com/schemas/customer",

  "type": "object",
  "properties": {
    "first_name": { "$ref": "#/definitions/name" },
    "last_name": { "$ref": "#/definitions/name" },
    "shipping_address": { "$ref": "/schemas/address" },
    "billing_address": { "$ref": "/schemas/address" }
  },
  "required": ["first_name", "last_name", "shipping_address", "billing_address"],

  "definitions": {
    "name": { "type": "string" }
  }
}
```

\$ref isn't just good for avoiding duplication. It can also be useful for writing schemas that are easier to read and maintain. Complex parts of the schema can be defined in definitions with descriptive names and referenced where it's needed. This allows readers of the schema to more quickly and easily understand the schema at a high level before diving into the more complex parts.

Note: It's possible to reference an external subschema, but generally you want to limit a \$ref to referencing either an external schema or an internal subschema defined in definitions.

5.5 Recursion

The \$ref keyword may be used to create recursive schemas that refer to themselves. For example, you might have a person schema that has an array of children, each of which are also person instances.

{ json schema }

```
{
  "type": "object",
  "properties": {
    "name": { "type": "string" },
    "children": {
      "type": "array",
      "items": { "$ref": "#" }
    }
  }
}
```

A snippet of the British royal family tree



```
{
  "name": "Elizabeth",
  "children": [
    {
      "name": "Charles",
      "children": [
        {
          "name": "William",
          "children": [
            { "name": "George" },
            { "name": "Charlotte" }
          ]
        },
        {
          "name": "Harry"
        }
      ]
    }
  ]
}
```

Above, we created a schema that refers to itself, effectively creating a “loop” in the validator, which is both allowed and useful. Note, however, that a `$ref` referring to another `$ref` could cause an infinite loop in the resolver, and is explicitly disallowed.

{ json schema }

```
{
  "definitions": {
    "alice": { "$ref": "#/definitions/bob" },
    "bob": { "$ref": "#/definitions/alice" }
  }
}
```

5.6 Bundling

Working with multiple schema documents is convenient for development, but it is often more convenient for distribution to bundle all of your schemas into a single schema document. This can be done using the `$id` keyword in a subschema. When `$id` is used in a subschema, it creates a new *Base URI* (page 73) that any references in that subschema and any descendant subschemas will resolve against. The new *Base URI* (page 73) is the value of `$id` resolved against the *Base URI* (page 73) of the schema it appears in.

Draft 4

In Draft 4, `$id` is just `id` (without the dollar sign).

This example shows the customer schema example and the address schema example bundled into a single schema document.

{ json schema }

```
{
  "$id": "https://example.com/schemas/customer",

  "type": "object",
  "properties": {
    "first_name": { "type": "string" },
    "last_name": { "type": "string" },
    "shipping_address": { "$ref": "/schemas/address" },
    "billing_address": { "$ref": "/schemas/address" }
  },
  "required": ["first_name", "last_name", "shipping_address", "billing_address"],

  "definitions": {
    "address": {
      "$id": "/schemas/address",

      "type": "object",
      "properties": {
        "street_address": { "type": "string" },
        "city": { "type": "string" },
        "state": { "$ref": "#/definitions/state" }
      },
      "required": ["street_address", "city", "state"],

      "definitions": {
        "state": { "enum": ["CA", "NY", "... etc ..." ] }
      }
    }
  }
}
```

Notice that the `$ref` keywords from the customer schema resolve the same way they did before except that the address schema is now defined at `/definitions/address` instead of a separate schema document. You should also see that `"$ref": "#/definitions/state"` resolves to the `definitions` keyword in the address schema rather than the one at the top level schema like it would if the subschema `$id` wasn't used.

You might notice that this creates a situation where there are multiple ways to identify a schema. Instead of referencing `/schemas/address` (<https://example.com/schemas/address>) You could have used `#/definitions/address` (<https://example.com/schemas/customer#/definitions/address>). While both of these will work, the one shown in the example is preferred.

Note: It is unusual to use `$id` in a subschema when developing schemas. It's generally best not to use this feature explicitly and use schema bundling tools to construct bundled schemas if such a thing is needed.

CHAPTER 6

Acknowledgments

Michael Droettboom wishes to thank the following contributors:

- Alexander Kjeldaas
- Alexander Lang
- Anders D. Johnson
- Armand Abric
- Ben Hutton
- Brandon Wright
- Brent Tubbs
- Chris Carpenter
- Christopher Mark Gore
- David Branner
- David Michael Karr
- David Worth
- E. M. Bray
- Fenhl
- forevermatt
- goldaxe
- Henry Andrews
- Hervé
- Hongwei
- Jesse Claven
- Koen Rouwhorst

- Mike Kubit
- Oliver Kurmis
- Sam Blackman
- Vincent Jacques

Index

Symbols

\$comment, 47
\$id, 72, 74, 78
\$ref, 75, 77
\$schema, 66, 67

A

additionalItems, 38, 40
additionalProperties, 32
allOf, 51
annotation, 46
anyOf, 51
array, 37
 contains, 42
 items, 38
 length, 43
 list validation, 38
 tuple validation, 39
 tuple validation; additionalItems, 40
 uniqueness, 43

B

base URI, 73
boolean, 44
bundling, 78

C

canonical URI, 74
comment, 47
conditionals, 56, 60
 dependencies, 57
 else, 60
 if, 60
 implication, 64
 property dependencies, 57
 schema dependencies, 59
 then, 60
const, 48
constant values, 48

contains, 38, 42
contentEncoding, 49
contentMediaType, 49

D

date, 20
date-time, 20
default, 46
definitions, 76
dependencies, 57
description, 46

E

else, 60
email, 21
enum, 47
enumerated values, 47
examples, 46
exclusiveMaximum, 26
exclusiveMinimum, 26

F

format, 20
 date, 20
 date-time, 20
 email, 21
 hostname, 21
 idn-email, 21
 idn-hostname, 21
 ipv4, 21
 ipv6, 21
 iri, 21
 iri-reference, 21
 json-pointer, 21
 regex, 22
 relative-json-pointer, 21
 time, 20
 uri, 21
 uri-reference, 21
 uri-template, 21

H

hostname, 21

I

idn-email, 21
idn-hostname, 21
if, 60
implication, 64
integer, 23
ipv4, 21
ipv6, 21
iri, 21
iri-reference, 21
items, 38

J

JSON Pointer, 72
json-pointer, 21

M

maximum, 26
maxItems, 43
maxLength, 18
maxProperties, 36
media, 49
 contentEncoding, 49
 contentMediaType, 49
minimum, 26
minItems, 43
minLength, 18
minProperties, 36
multipleOf, 25

N

named anchors, 72
non-JSON data, 49
not, 53
null, 45
number, 23
 multiple of, 25
 range, 26

O

object, 29
 properties, 30, 32
 properties; regular expression, 31
 property names, 36
 required properties, 34
 size, 36

oneOf, 52

P

pattern, 19
patternProperties, 31

properties, 30
property dependencies, 57
propertyNames, 36

R

readOnly, 46
recursion, 77
regex, 22
regular expressions, 22
relative-json-pointer, 21
required, 34
retrieval URI, 73

S

schema
 custom, 67
 extending, 67
 keyword, 66
schema composition, 50
 allOf, 51
 anyOf, 51
 not, 53
 oneOf, 52
 subschema independence, 53
schema dependencies, 59
schema identification, 71
string, 17
 format, 20
 length, 18
 regular expression, 19
structure, 69
structuring
 \$ref, 75
 base URI, 73
 base URI; \$id, 74
 base URI; retrieval URI, 73
 bundling; \$id, 78
 definitions, 76
 recursion; \$ref, 77
 schema identification, 71
 subschema identification; \$id, 72
 subschema identification; JSON Pointer, 72

T

then, 60
time, 20
title, 46
type, 15
types
 basic, 15
 numeric, 23

U

uniqueItems, 43

[uri](#), 21
[uri-reference](#), 21
[uri-template](#), 21

W

[writeOnly](#), 46